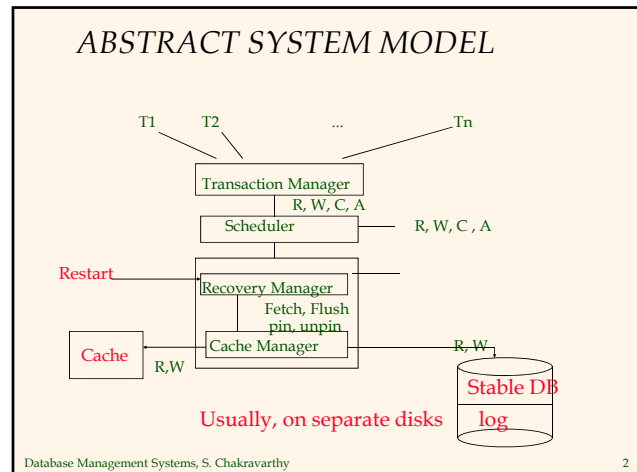


## Transaction Management Recovery Chapter 18 (3<sup>rd</sup> ed)

Instructor: Sharma Chakravarthy  
sharma@cse.uta.edu  
The University of Texas @ Arlington

Database Management Systems, S. Chakravarthy 1



### Comments

- ❖ Scheduler orders the operations so that the execution is **serializable and recoverable**
- ❖ The log contains records of the form [Ti, x, before\_val, after\_val] – Ti has written value after\_val for data item x which had value before\_val
- ❖ We deal with recovery when the **volatile storage fails (abort rollback is easier than recovery)**
  - Last committed value of x
  - Committed database state
  - Before (for undo) and after images (for redo) of x

Database Management Systems, S. Chakravarthy 3


### Recoverability

- ❖ When is recoverability required
  - **When an aborted Tx writes to stable storage (DB)**
    - ◆ If a system failure occurs at this point, the DB contains the effects of an aborted tx
    - ◆ These effects must be undone (UNDO)
  - **If a tx commits and does not write all of its updates to the DB**
    - ◆ This can happen if the buffer has not yet flushed the page to disk!!
    - ◆ If a system failure occurs at this point, the DB does NOT contain the effects of the committed tx
    - ◆ These effects must be redone (REDO)

Database Management Systems, S. Chakravarthy 4

## Motivation

- ❖ Atomicity:
  - Transactions may abort ("Rollback").
- ❖ Durability:
  - What if DBMS stops running? (Causes?)
- ❖ **Desired Behavior** after system restarts:
  - T1, T2 & T3 should be **durable**.
  - T4 & T5 should be **aborted** (effects not seen).




## Recovery Techniques

1. Incremental log with **deferred** updates (redo / no undo)
  - No writes before commit. Hence no undo. Redo is needed as committed Txs might not have written to disk
2. Incremental log with **immediate (in place)** updates (redo / undo)
  - If updates are happening in place, uncommitted Txs have written to disk by bm; hence undo.
  - If commit is not atomic, you need redo.
3. Shadow Paging (**alternative** to log-based crash recovery - no undo/no redo)
  - Pointers are swapped **atomically!!**

## Shadow paging

### Example of Shadow Paging

Shadow and current page tables after write to page 4



## Recovery

- ❖ **Tx UNDO**: removes all the effects of this Tx only
- ❖ **Global Undo**: when recovering from a system failure, the effects of all **incomplete** Txs have to be rolled back
- ❖ **Partial REDO**: when recovering from a system failure, results of committed Txs have to be redone (because some of the committed results are still in the buffer)
- ❖ **Global REDO**: archive recovery. Apply all committed Txs to a backup copy to bring it to current state.

## Assumptions

- ❖ Concurrency control is in effect.
  - **Strict 2PL**, in particular.
- ❖ Updates are happening “in place” (immediate)
  - i.e., data is overwritten on (deleted from) the disk.
  - As we have seen, this is done by the buffer manager and applications/TM do not have any control over it
- ❖ Looking for a simple scheme to guarantee Atomicity & Durability?

## Recoverability

- ❖ When is recoverability required
  - **When an aborted Tx writes to stable storage (DB)**
    - ◆ If a system failure occurs at this point, the DB contains the effects of an aborted tx
    - ◆ **These effects must be undone (UNDO)**
  - **If a tx commits and does not write all of its updates to the DB**
    - ◆ This can happen if the buffer has not yet flushed the page to disk!!
    - ◆ If a system failure occurs at this point, the DB does NOT contain the effects of the committed tx
    - ◆ **These effects must be redone (REDO)**

## Buffer Management and *undo*

- ❖ Replacement algorithms write dirty pages to disk. This is controlled **solely** by the buffer manager.
  - If no dirty pages are written back to disk, transaction undo can be limited to buffer operations
- ❖ Disadvantage: requires very large buffers.
- ❖ **Steal**: modified buffer pages can be written to disk (due to buffer replacement) at any time (**even before commit**)
- ❖ **No Steal**: modified pages are kept in buffer **at least until the end** of the transaction (EOT)
  - Unrealistic; requires large and varying size buffer space!

## Buffer Management and *Redo*

- ❖ When a tx commits, all its pages must be written to disk. Otherwise, durability has to be guaranteed in some other way (redo).
  - If all dirty pages are written at the end of Tx (still need to be atomic), no logging is required and hence no redo.
- ❖ Disadvantage: requires too many I/O's and waiting for EOT to finish writing.
- ❖ **Force**: All modified buffer pages are written at EOT
  - Commit has to wait for I/O to finish; **increases response time**
  - System failure can **still occur** during commit process
- ❖ **No Force**: No writing is triggered at the time of EOT (only decided by the BM)

## Handling the Buffer Pool

- ❖ **Force** writes to disk just before commit?

- Poor response time.
- But provides durability.
- **No redo**

- ❖ **Steal** buffer-pool frames from uncommitted Tx's?

- If not, need very large buffers
- If so, how can we ensure atomicity?
- **undo**

	No Steal	Steal
Force	Trivial (if commit is atomic)	Undo/ no redo
No Force	Redo/ no undo	Desired Need both Redo/undo

## More on Steal and Force

- ❖ **STEAL** (why enforcing Atomicity is hard)

- To steal frame F: Current page in F (say P) may be written to disk; some Xact may hold lock on P.
  - ♦ What if the Xact with the lock on P aborts? (locking is different from pinning)
  - ♦ Must remember the old value of P at steal time (to support UNDOing the write to page P).

- ❖ **NO FORCE** (why enforcing Durability is hard)

- What if system crashes before a modified page (of a committed Tx) is written to disk?
- Write as little as possible, in a convenient place (log), at commit time, to support REDOing modifications.

## Basic Idea: Logging

- ❖ Record REDO and UNDO information, for **every** update (insert/delete/modify), in a **log**.

- Sequential writes to log (put it on a **separate** disk).
- Minimal info (diff) written to log, so multiple updates fit in a single log page.

- ❖ **Log**: An ordered list of REDO/UNDO actions

- Log record contains:
  - <XID, pageID, offset, length, old value, new value>
- and additional control info (which we'll see soon).

## Write-Ahead Logging (WAL)

- ❖ The **Write-Ahead Logging** (WAL) Protocol:

- Do not flush an **uncommitted** update to the stable database (or storage) until the log record containing its before image has been flushed to the log.
  - ♦ Otherwise, you cannot undo it!!

- That is, must **force** the **log record** for an update **before** the corresponding **data page** gets to disk.

- ❖ Guarantees Atomicity (using undo or rollback)

## The Force-at-commit Rule

- ❖ The Force-at-Commit Protocol:
  - Do not commit a Tx until the after images of all its updated pages (as log) are in stable storage.
    - ◆ Otherwise, cannot do redo!!
  - That is, must **write/flush all log records** for a Xact **before commit** is declared!
- ❖ Guarantees Durability (by doing redo or roll forward)
- ❖ NOTE: only write all **log records** before commit, **not data pages!!**
  - Log information is much smaller!

## Info needed for undo and redo

- ❖ All the implementations observe the above two rules to ensure that **sufficient (and minimal)** information is stored in the log:
  - UNDO Rule (WAL) : log/store the value of x in a stable storage before overwriting it by an uncommitted value (store before image)
    - ◆ Each update is logged!
  - REDO Rule (Force-at-Commit) : Before a tx commits, the value it wrote for each data item must be logged/stored in a stable storage (store committed image).

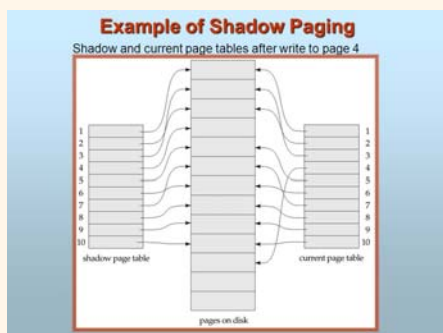
## Types of Recovery algorithms

- ❖ Both Undo/Redo
  - corresponds to steal/no force
- ❖ Undo/no-redo
  - corresponds steal/Force
- ❖ No-undo/redo
  - corresponds to no steal/No force
- ❖ No-undo/no-redo
  - corresponds to no steal/force
  - Shadow paging (used in early System R)
- ❖ Most/all commercial systems use both **undo/redo algorithms**

## Atomicity of Commit

- Flush all of T1's updates  
 Before commit to avoid redo  
 ↓  
 T1: start .... Write(p) ..... Commit (takes finite amount of time)  
 ↑  
 Flush all of T1's  
 Updates only after  
 Commit to avoid undo
- ❖ One of (Undo, Redo) **cannot be avoided** unless **commit is made atomic**. That is exactly what shadow paging did by making it an atomic pointers swap. Not realistic for large databases

## Shadow paging



## Restart

- ❖ Restart requires a fair bit of book keeping
  - It needs to know which txs were **active** at the time of failure - so it can abort them
  - It needs to know which updates of **committed Tx**s were **not written** to the stable db - so it can redo them
- ❖ Moreover, restart must be **fault-tolerant**. That is, if the system fails when restart is running, it must be possible to re-execute restart (as many times as needed) and get the correct result (DB state)
  - Hence the restart algorithm must be **idempotent**
  - property of certain operations in mathematics and computer science, that can be applied multiple times without changing the result beyond the initial application.
- ❖ Examples: **absolute function**, **multiplication by 1**, **max**
- ❖ **Non-idempotent operators**: **general addition**, **multiplication**

## Recovery characteristics

- ❖ Recovery should add **as little overhead** to normal processing as possible
  - Avoid excessive flushing (I/O)
  - Avoid logging too much data
- ❖ Recovery should be done **quickly**
  - So system is down for a short period
  - High availability
- ❖ **This is even more critical today due to 24/7 availability expectation!**

## Log Records

### LogRecord fields:

update records only {


- prevLSN
- XID
- type
- pageID
- length
- offset
- before-image
- after-image

### Possible log record types:

- ❖ **Update**
- ❖ **Commit**
- ❖ **Abort**
- ❖ **End** (signifies end of commit or abort)
- ❖ **Each log record has a unique, non-decreasing id**
- ❖ **Both WAL and Force-at-commit are being followed!**

### Recovery actions

- ❖ In the case of the DB being modified by incomplete Tx's (steal policy by BM), to what extent (how far back) does the log have to be processed for UNDO recovery?
- ❖ **Oldest incomplete transaction**




Database Management Systems, S. Chakravarthy

25

### Recovery actions

- ❖ If the DBMS does not use a **force** (no-force by BM) discipline, how much of the log (how far back) has to be processed for redo recovery?
- ❖ **Depends on when the buffer replaces the page**
- ❖ You can have a page in the buffer that has been modified by many committed and some uncommitted txs for a long time!



Database Management Systems, S. Chakravarthy

26

### Recovery actions or how much

- ❖ **Depends on when the buffer replaces the page**
- ❖ If there is a hot spot, the buffer page for that hot spot will contain changes of many committed transactions.
- ❖ Hence redo recovery will have to go back very far in the log (expensive)
- ❖ Depends on the interval between crashes
- ❖ Higher the availability of the system, the more costly recovery will become.
- ❖ **Hence, Need checkpoints: to limit the scope of redo**
- ❖ **Note that there is no way to limit scope of undo except to limit Tx to a certain size!**

Database Management Systems, S. Chakravarthy

27

### Checkpointing

- ❖ Any activity that is done during normal processing to reduce the amount of work to redo after recovery.
- ❖ Involves 3 steps
  - Write a begin\_checkpoint record to the log file
  - Write all checkpoint data to the log file (shall see soon)
  - Write an end\_checkpoint record to the log file
- ❖ **In fact, SQL allows you to checkpoint a transaction using the CHECKPOINT statement!**
- ❖ **All DBMSs do checkpointing periodically (even if not explicitly specified) to reduce redo**

Database Management Systems, S. Chakravarthy

28

## *Tx-oriented checkpointing*


- ❖ FORCE discipline can be seen as checkpointing at the Tx level
  - That is, writing all modified pages at commit!
- ❖ Disadvantages:
  - Large number of I/O's at EOT
  - Hotspots have to be written to disk very often.

## *Tx-consistent checkpointing*

- ❖ Global in nature
- ❖ Creates a transaction consistent database
- ❖ Requires all update activity be **quiescent**
  - All incomplete tx are completed and new ones are **not admitted**
  - The checkpoint is actually generated when the last update is completed
- ❖ After the end\_checkpoint record has been written, normal operation resumes
- ❖ Stop the pipeline temporarily, empty all Txs, and resume!
- ❖ **Not realistic!**

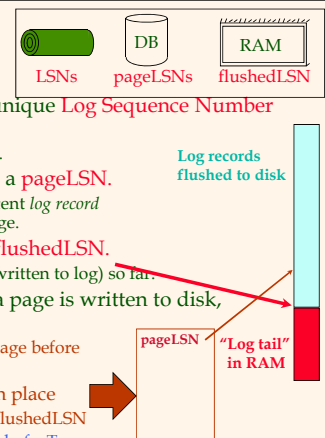
## *Cache consistent checkpointing*

- ❖ Also known as action-consistent checkpoint
- ❖ Stops processing any new operations (temporarily leave an active tx in blocked state)
- ❖ Flushes all the dirty pages in the cache/buffer
- ❖ Places markers at the end of the log and abort list to indicate that the flushes took place.
- ❖ Response time increases, not realistic!



## *WAL & the Log*

- ❖ Each log record has a unique **Log Sequence Number (LSN)**.
  - LSNs always increasing.
- ❖ Each **data page** contains a **pageLSN**.
  - The LSN of the most recent log record for an update to that page.
- ❖ System keeps track of **flushedLSN**.
  - The max LSN flushed (written to log) so far.
- ❖ **WAL protocol:** Before a page is written to disk,
  - pageLSN ≤ flushedLSN (log record is in stable storage before flushing the data page)
- ❖ **Force-at-commit** is also in place
  - lastLSN (of that Tx) ≤ flushedLSN
  - lastLSN: latest log record of a Tx





## Goals of ARIES\*

- ❖ ARIES is an “industrial strength” buffer management and logging/recovery scheme developed at IBM Almaden Research (by C. Mohan and others)
  - **Used by most commercial DBMSs today!**
  - no constraints on buffer fetch and eviction
    - ♦ *steal*
  - ♦ support for long-running transactions
  - fast commit
    - ♦ *no-force*
  - on-line incremental “fuzzy” checkpointing
    - ♦ fully concurrent with automatic log truncation
  - fast recovery, restartable (or idempotent) if the system fails while recovering

\*Algorithms for Recovery and Isolation Exploiting Semantics (or ARIES)

## Introduction to ARIES

- ❖ 1. Every log record is tagged with a monotonically increasing Log Sequence Number (LSN).
  - ♦ At recovery, log records can be retrieved efficiently by LSN.
- ❖ 2. Keep a **transaction table** in memory, with a record for each active transaction.
  - ♦ Keep each transaction's **lastLSN** of its most recent log record.
- ❖ 3. Maintain a backward-linked list (in the log) of log records for each transaction (**prevLSN**)
  - ♦ (Write the transaction's current **lastLSN** into each new log record.)
- ❖ 4. Each record in the log pertains to exactly one page, whose ID is logged as part of the record.
  - Means cannot have records spanning a page!

## Log Records

### LogRecord fields:

prevLSN  
 XID  
 type  
 pageID  
 length  
 offset  
 before-image  
 after-image

update  
 records  
 only

### Possible log record types:

- ❖ **Update**
- ❖ **Commit**
- ❖ **Abort**
- ❖ **End** (signifies end of commit or abort)
- ❖ **Compensation Log Records (CLRs)**
  - for UNDO actions

## Types of log records

1. **Updating a page**
  - After modifying a page, an update **record** with before and after image is appended to the log buffer
  - The **pageLSN** of the modified page is set to the LSN of the log record
  - The page must be **pinned before** modifying the page and **unpinned after**
2. **Commit**
  - This log record is written when a **commit is encountered/started (not finished)**
  - This and prev log records are **force written** to the log file (force-at-commit rule)
  - Note that commit is not complete until further book keeping actions (e.g., removing Tx from Tx table)

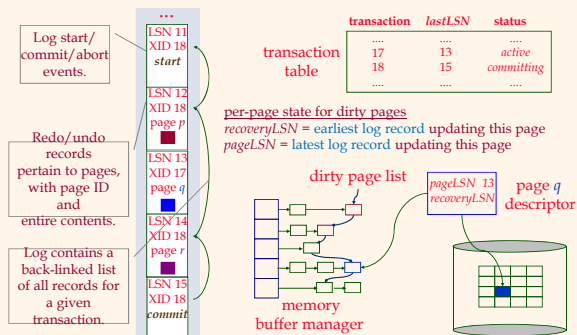
## Types of log records (2)

3. **Abort**
  - This log record is written when an abort is encountered and **undo is initiated**
  - Remember that abort takes **finite amount of time to finish**
4. **End:**
  - When a Tx is aborted and committed, some additional actions are necessary. After these additional actions are done, an **end type log record** is appended to the log
5. **Undoing an update:**
  - when a Tx is rolled back (either for abort or for rollback), its updates are undone. When the action described by an update record is undone, a CLR is written. **It will have only the before image for redoing it.**
    - ♦ CLR's are redone but never undone!

## Other Log-Related State (in main memory)

- ❖ **Transaction Table (TxT):**
  - One entry per active Xact.
  - Contains **XID**, **status** (running/committed/aborted), and **lastLSN** (points to the previous LSN of the **SAME** transaction)
- ❖ **Dirty Page Table (DPT):**
  - One entry per each dirty page in buffer pool.
  - Contains **recoveryLSN** - the LSN of the log record which **first** caused the page to be dirty (**not the same as dirty bit**)
- ❖ In addition, note that each page has **pageLSN** - the latest log record that modified that page
  - Useful when it is written to disk; we know that the log was flushed up to that **pageLSN** **before writing the page to disk (dirty page)**

## ARIES Data Structures



## Normal Execution of a Tx

- ❖ Series of **reads & writes**, followed by **commit** or **abort**.
  - We will assume that write is atomic on disk.
    - ♦ In practice, additional details need to be worked out to deal with non-atomic writes.
    - ♦ Some provide atomic writes, Writing and reading back to make sure it is correct!
- ❖ **Strict 2PL.**
- ❖ **STEAL, NO-FORCE** buffer management, with **Write-Ahead Logging and force-at-commit.**

## The Big Picture: What's Stored Where



### LogRecords

prevLSN  
XID  
type  
pageID  
length  
offset  
before-image  
after-image



**Data pages**  
each  
with a  
pageLSN

**master record**  
has Last chkpt  
Record lsn



### Xact Table

lastLSN  
status

**Dirty Page Table**  
recLSN

**flushedLSN**

**BM manages pages  
When in RAM**


## Checkpointing

- ❖ Periodically, the DBMS creates a **checkpoint**, in order to minimize the time taken to recover in the event of a system crash. Write to log:
  - **begin\_checkpoint** record: Indicates when chkpt began.
  - **end\_checkpoint** record: Contains current *Xact table* and *dirty page table*. This is a 'fuzzy checkpoint' because:
    - ◆ Other Xacts continue to run; so these tables are accurate only as of the time of the **begin\_checkpoint** record.
    - ◆ No attempt to force dirty pages to disk; effectiveness of checkpoint limited by oldest unwritten change to a dirty page. (So it's a good idea to periodically flush dirty pages to disk!)
  - After writing the **end\_checkpoint** record, Store LSN of begin chkpt record in a safe/known place (**master record**). Implication?
- ❖ Note the difference between fuzzy checkpointing and others discussed earlier
  - Normal operations continue during checkpointing

## Simple Transaction Abort

- ❖ For now, consider an explicit abort of a Xact.
  - No crash involved.
  - **Memory not lost!!**
- ❖ We want to "play back" the log in reverse order, UNDOing updates.
  - Before starting UNDO, write an **Abort log record** (not the end record)
    - ◆ For recovering from crash during UNDO!
  - Get **lastLSN** of Xact from **Xact table**.
  - Follow chain of log records backward via the **prevLSN** field (until null prevLSN is encountered)

## Abort, cont.

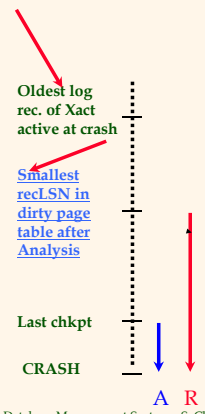


- ❖ To perform UNDO, must have a lock on **data record!**
  - No problem!
  - May need to bring that page from disk! To BM
  - Do you need a lock on the **log record** page?
  - Do you need to pin it?
- ❖ Before restoring old value of a page, write a CLR (has only before image): (we will use this later for redo)
  - You continue logging (CLRs) while you UNDO!!
  - CLRs prevLSN is interpreted as: **undonextLSN**
    - ◆ Points to the next LSN to undo (i.e., the prevLSN of the record we're currently undoing).
  - CLRs **never** Undone (but they might be Redone when repeating history: guarantees Atomicity!)
- ❖ At end of UNDO, write an "end abort" log record

## Transaction Commit

- ❖ Write **commit** record to log. // not to disk
- ❖ All log records up to Xact's **lastLSN** are flushed.
  - Guarantees that **flushedLSN**  $\geq$  **lastLSN**.
  - Note that log flushes are sequential, synchronous writes to disk.
  - Many log records per log page.
- ❖ Commit() returns. // commit record is also flushed
- ❖ Write **end commit** record to log.
  - End record need not be flushed!
  - But flushing may save abort of a almost committed Tx
  - Correctness is not affected by not flushing!

## Crash Recovery: Big Picture

- 
- ❖ Start from a **checkpoint** (found via **master** record).
  - ❖ Three phases. Need to:
    - Figure out which Xacts have committed since checkpoint, which Xacts have failed (**Analysis**).
    - **REDO** *all* actions.
      - ♦ (repeat history)
    - **UNDO** effects of failed Xacts.
  - ❖ **Why redo before undo?**
- Relative positions of A, R, and U  
May be different than shown

## Recovery Phases

- ❖ **Restart-Analysis phase**: identifies dirty pages in the buffer pool (i.e., changes that have not been written to disk) and active Tx's at the time of crash
  - Starts with the last checkpoint record
  - There may be log recs between start and end checkpoint
  - Analysis phase may write "end" records if missing for aborted Tx's!
- ❖ **Redo**: Repeat all actions, starting from an **appropriate point in the log**, and restore the db state to what it was at the time of crash
  - Can be skipped if no pages in dirty page table!
- ❖ **Undo**: Undoes the actions of transactions that did not commit, so that the db reflects only the actions of committed transactions. **Also writes CLR's**
  - Skips logs based on CLR's; CLR's have been handled by redo!

## Log Records

### LogRecord fields:

update records only { prevLSN  
XID  
type  
pageID  
length  
offset  
before-image  
after-image

### Possible log record types:

- ❖ **Update**
- ❖ **Commit**
- ❖ **Abort**
- ❖ **End** (signifies end of commit or abort)
- ❖ **CLR (Compensation Log Records)**
  - for UNDO actions

## Example of a log

LSN LOG

00 update: T1 writes P5  
05 update: T2 writes P3  
10 T2 commit  
20 T2 End  
30 update: T3 writes P1  
40 update: T3 writes P3  
CRASH

Note: we do not write  
A begin tx type log record

T1 and T3 were active at the time of crash. Hence need to be undone; T2 is a committed Tx, and all its actions need to be written to disk; and P1, P3, and P5 are potentially dirty pages

## Recovery: The Analysis Phase

- ❖ Reconstruct state at checkpoint via **end\_checkpoint** record. That is,
  - From where to **do Redo**?
  - Dirty pages at the time of crash
  - Txs to **be undone** (those active at the time of crash)
- ❖ Scan log forward from checkpoint.
  - **End** record: Remove Xact from Xact table (if there)
  - **Other records (c, a, u, clr)**: Add Xact to Xact table,
    - ◆ set lastLSN=LSN, change Xact status to C on commit, else u
  - **Redoable (Update or CLR) record**:
    - ◆ If P **not** in Dirty Page Table (do nothing if it is in DPT)
      - Add P to DPT., set its recLSN=LSN (because we are coming from the ckpoint side; done only once!)
    - ◆ P **not** in DPT means that it was brought into memory after the checkpoint was written
- ❖ At the end of the Analysis phase, we have reconstructed Tx table and DPT in memory!

## Example of a log

DPT

Page	recLSN
P5	00
P3	05
P1	30

Tx Table

Tx	lastLSN	status
T1	00	u
T2	05, 10	c X
T3	30, 40	u

LSN LOG

00 update: T1 writes P5  
05 update: T2 writes P3  
10 T2 commit  
20 T2 End  
30 update: T3 writes P1  
40 update: T3 writes P3  
CRASH

Txt and DPT are assumed  
empty at checkpoint

T1 and T3 were active at the time of crash. Hence need to be undone; T2 as a committed Tx, and all its actions need to be written to disk; and P1, P3, and P5 are potentially dirty pages

## Recovery: The REDO Phase

recLSN = earliest log record updating this page (present only in the main memory DPT table)  
pageLSN = latest log record updating this page (present in every page in disk/buffer)

- ❖ We **repeat History** to reconstruct state at crash:
  - Reapply **all** updates (even of aborted Xacts!), redo CLR's.
- ❖ Scan forward from log rec containing **smallest recLSN** in DPT. For each **CLR or update** log rec **LSN**, REDO the action **unless**:
  1. Affected page is not in the Dirty Page Table, or
  2. Affected page is in DPT (and BM) but has **recLSN > LSN**, or
  3. **(in DPT and not in BM)**
    1. **pageLSN (in DB) ≥ LSN.** (page must be retrieved from DB)
- ❖ To **REDO** an action:
  - Reapply logged action.
  - Set **pageLSN** to **LSN of the log record**. No additional logging!
- ❖ At the **end of REDO phase**, end type records are written for all Txs with status C, and are removed from the Tx table. All committed Txs have been restored! **Why is this true?**

## Recovery: The UNDO Phase

ToUndo={ *l* | *l* is the lastLSN of a "loser" Xact }  
 //from the Tx table constructed

### Repeat:

- Choose largest LSN among ToUndo.
- If this LSN is a CLR and undonextLSN==NULL
  - ◆ Write an End record for this Xact.
- If this LSN is a CLR, and undonextLSN != NULL
  - ◆ Add undonextLSN to ToUndo
- Else this LSN is an update. Undo the update, write a CLR, add prevLSN to ToUndo.

Until ToUndo is empty.





Fig. 13. Restart recovery example with ARIES.

## Example of Recovery



Xact Table  
 lastLSN  
 status  
 Dirty Page Table  
 recLSN  
 flushedLSN

ToUndo


### LSN LOG

00 begin\_checkpoint  
 05 end\_checkpoint  
 10 update: T1 writes P5  
 20 update: T2 writes P3  
 30 T1 abort  
 40 CLR: Undo T1 LSN 10 NULL  
 45 T1 End  
 50 update: T3 writes P1  
 60 update: T2 writes P5  
 CRASH, RESTART

Page	recLSN
P5	10
P3	20
P1	50

Tx	lastLSN
T1	20
T2	60
T3	50

## Example: Crash During Restart!



Xact Table  
 lastLSN  
 status  
 Dirty Page Table  
 recLSN  
 flushedLSN

ToUndo

### LSN LOG

00,05 begin\_checkpoint, end\_checkpoint  
 10 update: T1 writes P5  
 20 update: T2 writes P3  
 30 T1 abort  
 40,45 CLR: Undo T1 LSN 10, T1 End  
 50 update: T3 writes P1  
 60 update: T2 writes P5  
 CRASH, RESTART  
 70 CLR: Undo T2 LSN 60  
 80,85 CLR: Undo T3 LSN 50, T3 end  
 CRASH, RESTART  
 90,95 CLR: Undo T2 LSN 20, T2 end

Page	recLSN
P5	10
P3	20
P1	50

Tx	lastLSN
T1	20
T2	70
T3	50

## Evaluating ARIES

- ❖ The ARIES logging/recovery algorithm has several advantages over other approaches:
  - *steal/no-force* with few constraints on buffer management
    - ◆ Steals act as incremental, nonintrusive checkpoints.
  - synchronous “fuzzy” checkpoints are fast and nonintrusive
  - minimizes recovery work
    - ◆ makes forward progress in failures during recovery
  - *repeating history* redo supports logical undo logging and alternative locking strategies (e.g., fine-grained locking)
- ❖ But: ARIES requires WAL with undos, LSNs written with every page, and redo records restricted to a single page.

## Additional Crash Issues (idempotency)

- ❖ What happens if system crashes
  - During Analysis?**
    - nothing has changed; do analysis again!
  - During REDO?**
    - things might have changed on disk; But we are going to REDO and restore it to the correct consistent state!!
  - During UNDO**
    - Whatever undo has been done they are NOT repeated again. Instead, REDO takes care of that. A successful undo is done only once!!!
- ❖ **How do you limit the amount of work in REDO?**
  - Flush asynchronously in the background.
  - Watch “hot spots”!
- ❖ **How do you limit the amount of work in UNDO?**
  - Cannot, except avoiding long-running Xacts.

## Summary of Logging/Recovery

- ❖ **Recovery Manager** guarantees Atomicity & Durability.
- ❖ Use WAL to allow STEAL/NO-FORCE without sacrificing correctness.
- ❖ LSNs identify log records; linked backwards for each Tx (via prevLSN).
- ❖ pageLSN allows comparison of data page and log records.

## Summary, Cont.

- ❖ **Checkpointing**: A quick way to limit the amount of log to scan on recovery.
- ❖ Recovery works in 3 phases:
  - **Analysis**: Forward from checkpoint.
  - **Redo**: Forward from oldest recLSN.
  - **Undo**: Backward from end to first LSN of oldest Xact alive at crash.
- ❖ While performing Undo, write CLR.
- ❖ Redo “repeats history”: Simplifies the logic!

Thank You !

